

Laptop Price Amongst Varying Suppliers

Arthur Starodynov (7751472)

June 07, 2022

Contents

Intro	1
Data	1
Visualization	5
Modeling	7
Model Performance	13
Conclusion	19

Intro

Intro

The purpose of this report is to figure out which brand of laptops provides the best laptop for the cheapest price.

Why might our predictive model be helpful?

The reason behind making our model useful and helpful is to guide customers to make an educated decision on the best laptop for their specific needs. Some customers might be intrigued and unsatisfied when a predictive model can show that some companies are not showing the fine print or real performance of the laptop.

```
#Read in the data
```

```
laptop.data <- read.csv('C:/Users/arthu/Dropbox/My PC (DESKTOP-9BV8I37)/Documents/PSTAT131Final/Cleaned  
set.seed(46)
```

Data

Introduction into the data set

As we all may know there are various laptop brands out there, with some of the most popular being Dell, Mac, and Lenovo. However, some of these popular laptop brands can be cheaping out customers based on their reputation that they have already built up. For example due to Apple's brand recognition and amazing customer support, the best performance laptop is mainly focused on performance features such as RAM, the Graphics Card, a Solid State Drive, and other factors. In addition, a factor that may persuade customers to

buy a specific laptop would be based on online reviews and generic ratings that professionals may include. Our goal through using this data set found on Kaggle, is to be able to figure out which laptop brand really bring the most high performing laptop with the best price. The data set below has 896 observations and 23 variables. Just by looking at the column names all of them are self explanatory where they state what is being described in each column.

Here we looked at what the values of our columns are within our data set.

```
laptop.data$ram_gb<-gsub(" GB GB","",as.character(laptop.data$ram_gb))
laptop.data$ssd<-gsub(" GB","",as.character(laptop.data$ssd))
laptop.data$hdd<-gsub(" GB","",as.character(laptop.data$hdd))
laptop.data$os_bit<-gsub("-bit","",as.character(laptop.data$os_bit))

print("Table view of all categorical variables")
```

```
## [1] "Table view of all categorical variables"
```

```
## [1] "Table view for column: brand"
```

```
##
##      acer ALIENWARE      APPLE      ASUS      Avita      DELL      HP      iball
##      58      4      28      254      18      154      142      1
##  Infinix  lenovo  Lenovo      LG      Mi MICROSOFT  MSI      Nokia
##      4      3      148      5      2      3      52      4
##  realme RedmiBook  SAMSUNG  Smartron      Vaio
##      4      3      1      3      5
```

```
## [1] "Table view for column: model"
```

```
##
##      14a      14s      15 15-ec1105AX      15q      15s
##      1      5      3      1      3      12
##      250      250-G6      3000      3511      430      A6-9225
##      1      1      1      1      1      1
##      Alpha      AMD      APU      Aspire      Asus      ASUS
##      1      1      3      24      1      13
##      Athlon      B50-70      Book  Book(Slim)      Bravo      Celeron
##      1      1      2      2      1      3
##  Chromebook  Commercial  CompBook  ConceptD      Cosmos      Creator
##      13      1      1      1      1      1
##      DA      DELL      Delta      Dual      E      EeeBook
##      1      2      1      1      4      4
##      Envy  ExpertBook      Extensa      F17      G15      G3
##      6      16      1      1      3      1
##      G5      G7      Galaxy      GAMING      GE76      GF63
##      2      1      1      1      1      6
##      GF65      GP65      GP76      Gram      GS      GS66
##      2      1      1      5      1      1
##      HP      Ideapad      IdeaPad      IDEAPAD      INBook      Inpiron
##      6      32      37      1      3      1
##  Inspiron  INSPIRON  Inspiron      Intel      Katana      Legion
##      77      2      1      4      4      11
##      Lenovo      Liber      MacBook      Missing      Modern      Nitro
##      1      13      28      95      10      5
##  Notebook      Omen      OMEN      Pavilion      Pentium      Predator
##      3      2      5      38      5      9
##  Prestige      Pro      Pulse      PURA      PureBook      Rog
```

```

##          5          4          3          4          4          1
##        ROG        Ryzen        SE        Spectre        Spin        Stealth
##         31         35         1         12         1         2
##       Summit      Surface      Swift      Sword      t.book      Thinkbook
##         1         3         8         2         3         4
##   ThinkBook   Thinkpad   ThinkPad   Thinpad   Travelmate      TUF
##         4         1        10         1         3        10
##        v15       V15       Vivo   VivoBook   VivoBook14      Vostro
##         1         1         1         89         1        33
##       WF65       X1       x360       X390       XPS       Yoga
##         1         1         2         1         5        14
##   Zenbook   ZenBook   Zephyrus
##         7        22         5
## [1] "Table view for column: processor_brand"
##
##      AMD      Intel      M1 MediaTek Qualcomm
##     208      660      24       3       1
## [1] "Table view for column: processor_name"
##
## A6-9225 Processor      APU Dual      Athlon Dual      Celeron Dual
##              1              7              2              24
##              Core      Core i3      Core i5      Core i7
##              1              170              312              112
##              Core i9      Core m3      Dual Core      Ever Screenpad
##              8              1              3              2
##      GeForce GTX      GeForce RTX      GEFORCE RTX      Genuine Windows
##              2              4              1              3
##      Hexa Core      M1 MediaTek Kompanio      Pentium Quad
##              2              24              3              14
##      Pentium Silver      Quad      Ryzen      Ryzen 3
##              2              1              1              26
##              Ryzen 5      Ryzen 7      Ryzen 9      Snapdragon 7c
##              85              58              26              1
## [1] "Table view for column: processor_gnrtn"
##
##    10th    11th    12th    4th    7th    8th    9th Missing
##    246    346     3     1    12    43     6     239
## [1] "Table view for column: ram_gb"
##
## 16 32 4 8
## 180 3 259 454
## [1] "Table view for column: ram_type"
##
##   DDR3   DDR4   DDR5  LPDDR3  LPDDR4  LPDDR4X
##    12   760    8    14    36    66
## [1] "Table view for column: ssd"
##
## 0 1024 128 2048 256 3072 32 512
## 151 111 12 2 201 1 1 417
## [1] "Table view for column: hdd"
##
## 0 1024 2048 512
## 666 164 1 65
## [1] "Table view for column: os"

```

```

##
##      DOS      Mac Windows
##      36      28      832
## [1] "Table view for column:  os_bit"
##
##  32  64
## 135 761
## [1] "Table view for column:  graphic_card_gb"
##
##   0   2   4   6   8
## 631 69 138 40  18
## [1] "Table view for column:  weight"
##
##      Casual      Gaming ThinNlight
##      566          39          291
## [1] "Table view for column:  display_size"
##
##   12.2    13    13.3    13.4    14    14.1    14.2    14.9    14.96    15
##     2     4     40     1    131     6     3     1     7     3
##   15.6    16    16.1    16.2    17.3 Missing
##   218    135     1     3     9    332
## [1] "Table view for column:  warranty"
##
##   0   1   2   3
## 332 521 30  13
## [1] "Table view for column:  Touchscreen"
##
## No Yes
## 793 103
## [1] "Table view for column:  msoffice"
##
## No Yes
## 606 290

##      brand      model processor_brand processor_name processor_gnrtn
##      0          95          0          0          239
##      ram_gb      ram_type          ssd          hdd          os
##      0          0          0          0          0
##      os_bit graphic_card_gb          weight    display_size    warranty
##      0          0          0          332          0
##      Touchscreen    msoffice    latest_price    old_price    discount
##      0          0          0          0          0
##      star_rating    ratings          reviews
##      0          0          0

##      brand      model processor_brand processor_name processor_gnrtn
##      0          95          0          0          239
##      ram_gb      ram_type          ssd          hdd          os
##      0          0          0          0          0
##      os_bit graphic_card_gb          weight    display_size    warranty
##      0          0          0          332          0
##      Touchscreen    msoffice    latest_price    old_price    discount
##      0          0          0          0          0
##      star_rating    ratings          reviews
##      0          0          0

```


```
##          0          0          0
```

Visualization

First Visualization


We first look at our target variable, Latest Price, convert it into USD \$ to make it easier to visualize and notice that the variable is right-skewed. Hence we apply a log transform and make a new target variable which is the log_transform of latest price.

```
laptop.data$latest_price <- laptop.data$latest_price*0.013
laptop.data$old_price <- laptop.data$old_price*0.013
d <- density(laptop.data$latest_price)
plot(d, main="Kernel Density of Latest Price")
polygon(d, col="red", border="blue")
```



```
laptop.data$log_latest_price <- log(laptop.data$latest_price)
d <- density(laptop.data$log_latest_price)
plot(d, main="Kernel Density of Log Latest Price")
polygon(d, col="red", border="blue")
```

Kernel Density of Log Latest Price



Cleaning of the data

We noticed that there were a few variables within our data set that had missing observations, so we needed to create a new column that would indicate if display size was given, and gave it either a “Yes or a “No value. In addition we wanted to use all the available data we had so any missing records were replaced with an average of the column, which were eventually grouped to make some levels, furthering the easiness of modeling later. In the clean data, we wanted to drop many of the columns that included N/A which were seen as `processor_gnrtn`, `discount`, `model`, `old_price`. Also we wanted to only maintain the same currency so we removed the `discount`, and `old_price` columns to not get clustered. With all the cleaning we made a new dataframe.

```
laptop.data["display_size_given"] <- laptop.data$display_size
laptop.data$display_size_given[!is.na(laptop.data$display_size_given)] <- "Yes"
laptop.data$display_size_given[is.na(laptop.data$display_size_given)] <- "No"
laptop.data$display_size <- as.numeric(laptop.data$display_size)
laptop.data$display_size <- round(na.aggregate(laptop.data$display_size),2)

laptop.data$brand <- tolower(laptop.data$brand)

laptop.data$brand <- ifelse(laptop.data$brand %in% c("acer",
                                                    "apple", "asus", "dell",
                                                    "hp","lenovo", "msi"),
                           laptop.data$brand,"other")
```

```

laptop.data$processor_brand <- tolower(laptop.data$processor_brand)

laptop.data$processor_brand <- ifelse(laptop.data$processor_brand %in% c("amd",
                                "intel"),
                                laptop.data$processor_brand, "other")

laptop.data$ram_gb_cat <- ifelse(as.numeric(laptop.data$ram_gb)<=8,
                                "less_than_8", "greater_than_8")

laptop.data$ssd<-as.numeric(laptop.data$ssd)
laptop.data$ssd_cat <- ifelse(laptop.data$ssd < 1024.0,
                                "less_than_1gb", "greater_than_1gb")

laptop.data$hdd<-as.numeric(laptop.data$hdd)
laptop.data$hdd_cat <- ifelse(laptop.data$hdd < 1024.0,
                                "low_hdd", "high_hdd")

laptop.data$processor_name <- tolower(laptop.data$processor_name)

laptop.data$processor_name <- ifelse(laptop.data$processor_name %in% c("celeron dual",
                                "core i3", "core i5", "core i7",
                                "m1", "pentium quad",
                                "ryzen 3", "ryzen 5",
                                "ryzen 7", "ryzen 9"),
                                laptop.data$processor_name, "other")

laptop.data.clean <-laptop.data[c(1,3,4,7,10,11,12,13,14,15,16,17,21,22,23,24,25,26,27,28)]

```

Modeling

Clustering of Data

Instead of using k folding technique we thought it would be more beneficial to be able to analyze using clustering analysis on the data set. Looking at the data set, we know that RAM, SSD, and the graphic card memory are performance attributes that customers will consider when they are in the market of buying a laptop. Hence, we decided to cluster upon these three columns.

```


tmp.laptop.data.clean <- laptop.data[,c(6,8,12)]
tmp.laptop.data.clean$ram_gb<- as.numeric(tmp.laptop.data.clean$ram_gb)
tmp.laptop.data.clean$ssd<- as.numeric(tmp.laptop.data.clean$ssd)
tmp.laptop.data.clean$graphic_card_gb<- as.numeric(
    tmp.laptop.data.clean$graphic_card_gb)

tmp.laptop.data.clean <- scale(tmp.laptop.data.clean)
rownames(tmp.laptop.data.clean) <- paste(laptop.data.clean$brand, "",
                                c(1:length(laptop.data.clean$brand)), sep="")

k1 <- kmeans(tmp.laptop.data.clean, centers = 2, nstart = 25)
fviz_cluster(k1, data = tmp.laptop.data.clean, labelsize=6)


```

Cluster plot

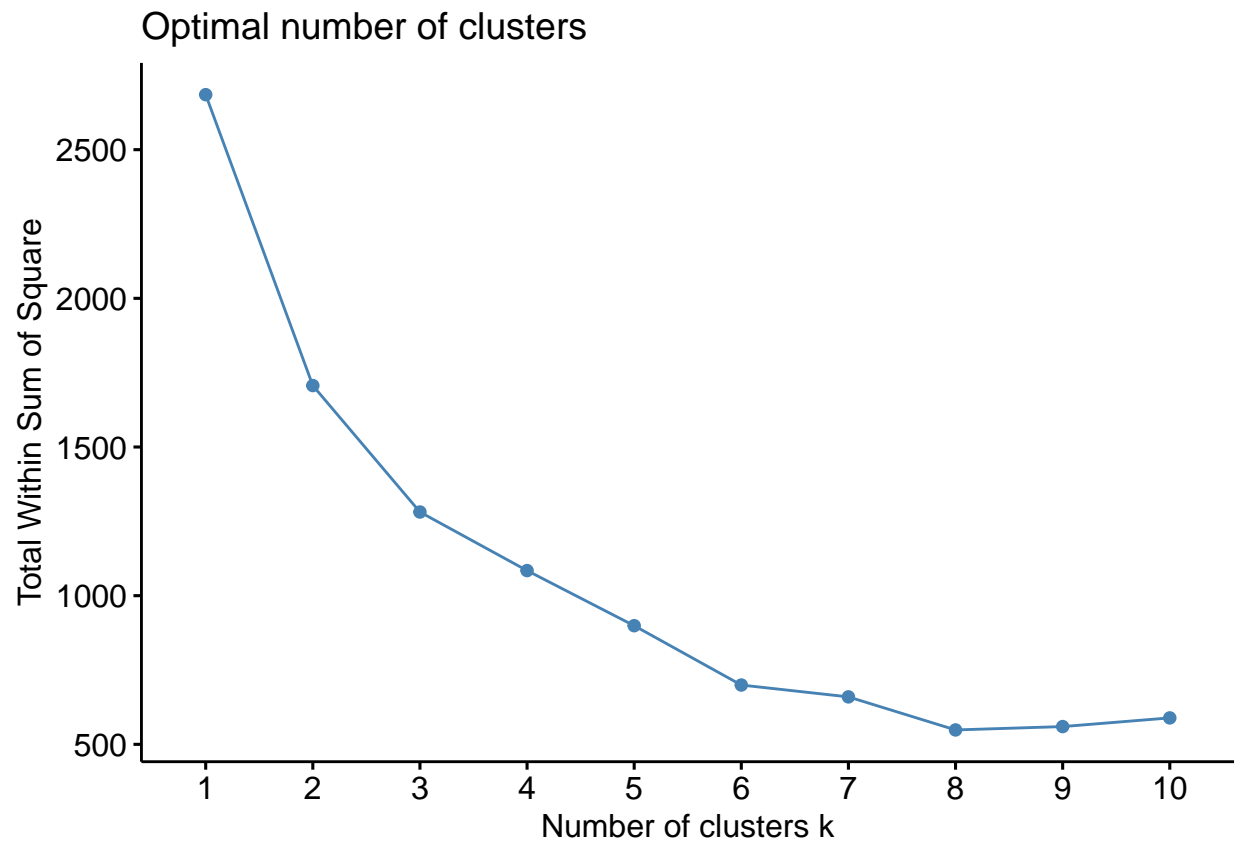


```
k2 <- kmeans(tmp.laptop.data.clean, centers = 4, nstart = 25)
fviz_cluster(k2, data = tmp.laptop.data.clean, labels=6)
```

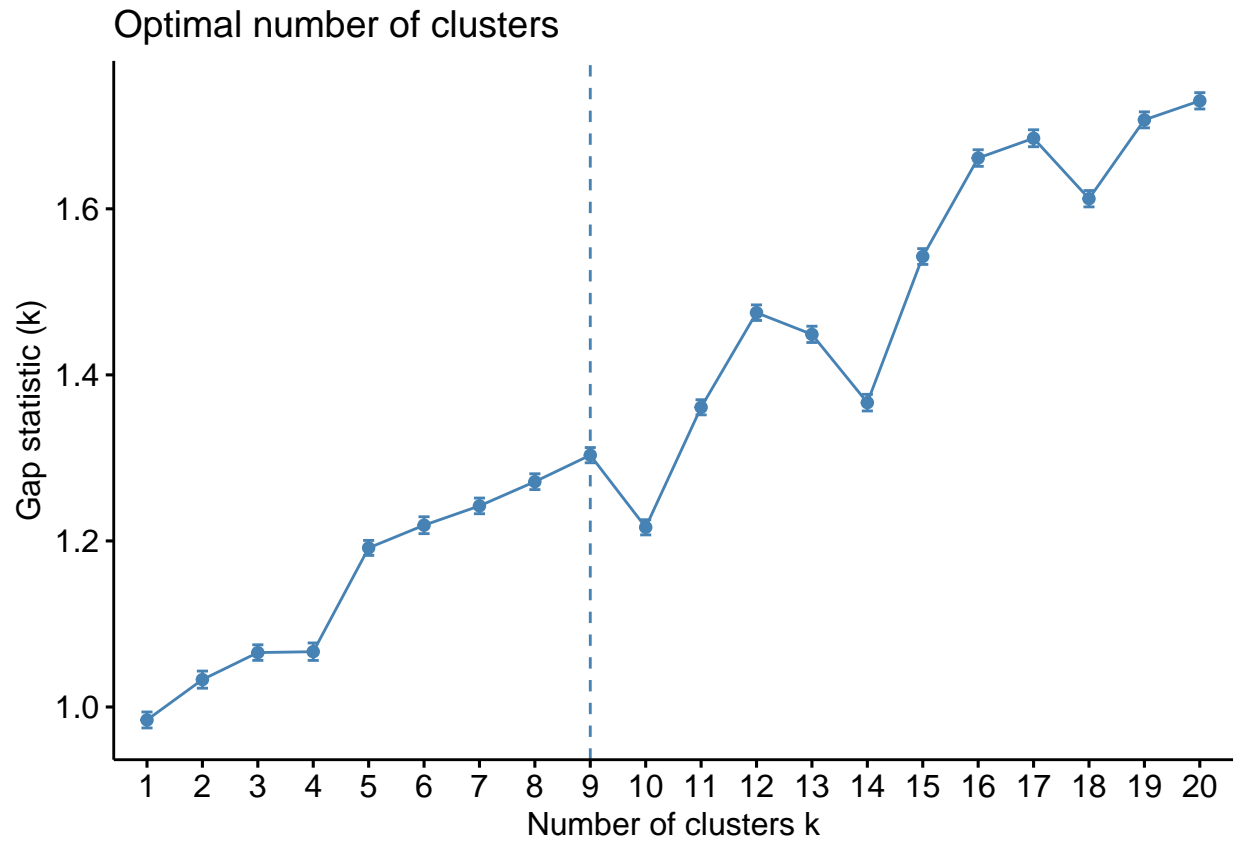

Cluster plot



```
fviz_nbclust(tmp.laptop.data.clean, kmeans, method = "wss")
```



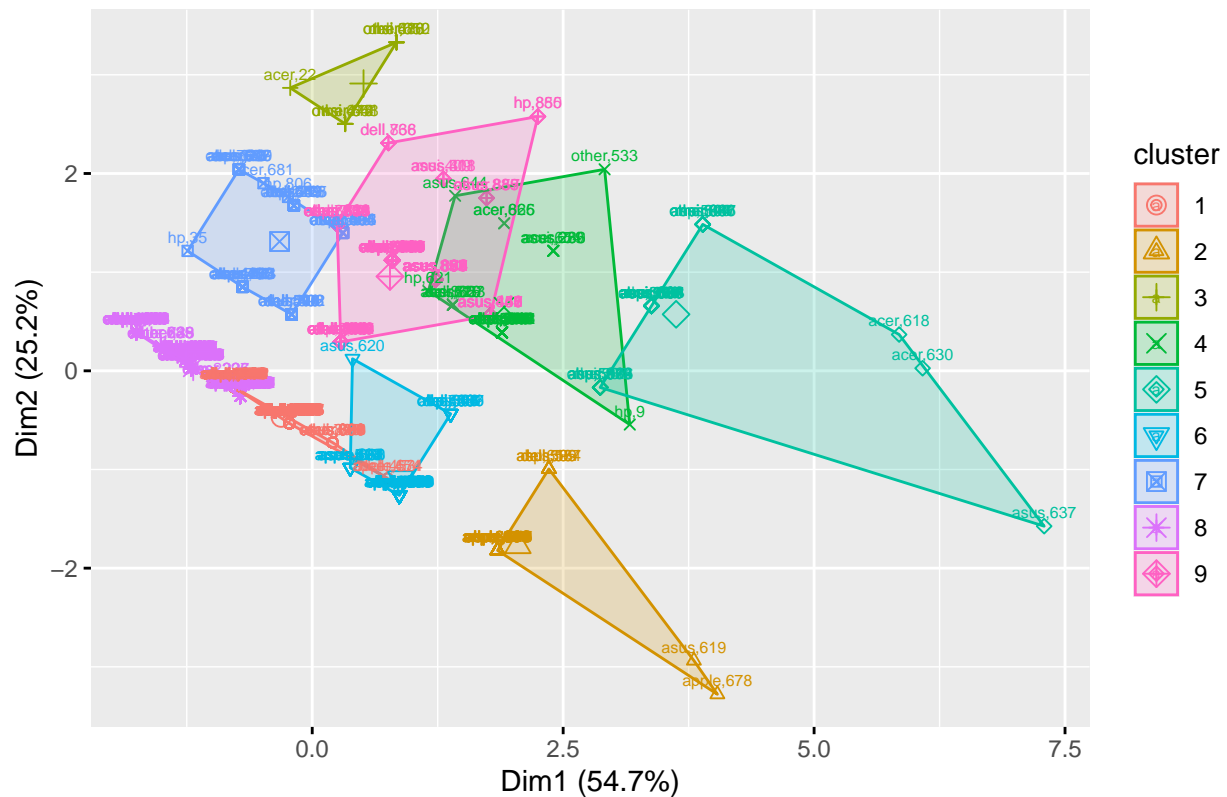
```
gap_stat <- clusGap(tmp.laptop.data.clean, FUN = kmeans, nstart = 25,  
                    K.max = 20, B = 50)  
fviz_gap_stat(gap_stat)
```



From the above results we notice that we were successful in clustering our laptop variables into separable clusters. Using a classic clustering method to find the optimal amount of them known as the elbow method, we figured out that the optimal amount of clusters is 9, so using 9 centers and the k means method we visualized these variables.

```
k3 <- kmeans(tmp.laptop.data.clean, centers = 9, nstart = 25)
fviz_cluster(k3, data = tmp.laptop.data.clean, labelsize=6)
```

Cluster plot



Below we printed the number of observations in each cluster which we will use to further visualize laptop price.

```
laptop.data.clean['cluster'] <- as.factor( k3$cluster)
table(laptop.data.clean$cluster)
```

```
##
##  1  2  3  4  5  6  7  8  9
## 256 33 11 37 37 76 51 283 112
```

Training the Data

Here we wanted to explore the data further so we implemented the technique to be able to fix the data into a training and testing data set.

```
train.indices <- sample(nrow(laptop.data.clean), floor(nrow(laptop.data.clean)/1.5), replace = FALSE)
validation.indices <- seq(nrow(laptop.data.clean))[-train.indices]
pred.laptop.train <- laptop.data.clean[train.indices,]
pred.laptop.train <- pred.laptop.train[,c(16,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,17,18,19,20,21)]
pred.laptop.validation <- laptop.data.clean[validation.indices,]
pred.laptop.validation <- pred.laptop.validation[,c(16,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,17,18,19,20,21)]
```

Model Performance

Analysis


Below we create our “recipe” or formula which we will use to fit our training data set. We wanted to use lasso regression to be able make the model have fewer features. We also applied a stepwise regression approach onto our model and compared the two.

```
glmnet.formula <- as.formula(log_latest_price ~ .)
glmnet.design.matrix <- model.matrix(glmnet.formula, data = pred.laptop.train)
dim(glmnet.design.matrix)
```

```
## [1] 597 49
```

```
glmnet.cv.laptop.out <- cv.glmnet(glmnet.design.matrix,
  y = pred.laptop.train$log_latest_price,
  family = c("gaussian"),
  type.measure="mse", # the model selection criteria
  alpha = 1) # The Lasso regression
```

```
plot(glmnet.cv.laptop.out)
```



```
saved.coef <- coef(glmnet.cv.laptop.out, s=c("lambda.1se"))
```

```
dim(saved.coef)
```

```
## [1] 50 1
```

```
chosen.vars <- data.frame(name = saved.coef@Dimnames[[1]][saved.coef@i + 1],  
                          coefficient = saved.coef@x)
```

```
print(paste("The lasso regression chose", dim(chosen.vars)[1]-1,  
            "variables and 1 intercept"))
```

```
## [1] "The lasso regression chose 29 variables and 1 intercept"
```

```
print(saved.coef)
```

```
## 50 x 1 sparse Matrix of class "dgCMatrix"  
##                               s1  
## (Intercept)                   6.7416819323  
## (Intercept)                   .  
## brandapple                    0.5558859411  
## brandasus                     -0.0020370338  
## branddell                     0.0048670322  
## brandhp                       .  
## brandmsi                      .  
## brandother                    .  
## processor_brandintel          .  
## processor_brandother          .  
## processor_namecore i3        -0.0097100993  
## processor_namecore i5         0.1985940376  
## processor_namecore i7         0.2869211681  
## processor_namei1              0.1365601575  
## processor_nameother           .  
## processor_namepentium quad -0.2503282289  
## processor_nameryzen 3        -0.0362703686  
## processor_nameryzen 5         0.0907451217  
## processor_nameryzen 7         0.1362783961  
## processor_nameryzen 9         0.3215072423  
## ram_typeDDR4                  .  
## ram_typeDDR5                  0.3005891972  
## ram_typeLPDDR3                0.3802361275  
## ram_typeLPDDR4                .  
## ram_typeLPDDR4X               .  
## osMac                        0.0012034446  
## osWindows                    -0.2671547136  
## os_bit64                      .  
## graphic_card_gb               0.0448758096  
## weightGaming                  .  
## weightThinNlight              0.0083628155  
## display_size                  0.0332215507  
## warranty                      0.0263151957
```

```
## TouchscreenYes          0.2750043592
## msofficeYes              .
## star_rating             -0.0281194638
## ratings                 -0.0000391211
## reviews                 .
## display_size_givenYes   .
## ram_gb_catless_than_8   -0.1557618406
## ssd_catless_than_1gb    -0.2561305868
## hdd_catlow_hdd          .
## cluster2                .
## cluster3                0.5303553466
## cluster4                .
## cluster5                0.1315500117
## cluster6                0.0068286880
## cluster7                .
## cluster8                -0.2059538924
## cluster9                .
```

```
glmnet.formula2 <- as.formula(log_latest_price ~ .)
glmnet.design.matrix.validation <- model.matrix(glmnet.formula2,
                                                data = pred.laptop.validation)

validation.preds.regularizedreg <- predict(glmnet.cv.laptop.out,
                                           newx = glmnet.design.matrix.validation,
                                           type = c("response"))
```

We can clearly see that the lasso model was fairly accurate, the stepwise regression model was still superior.

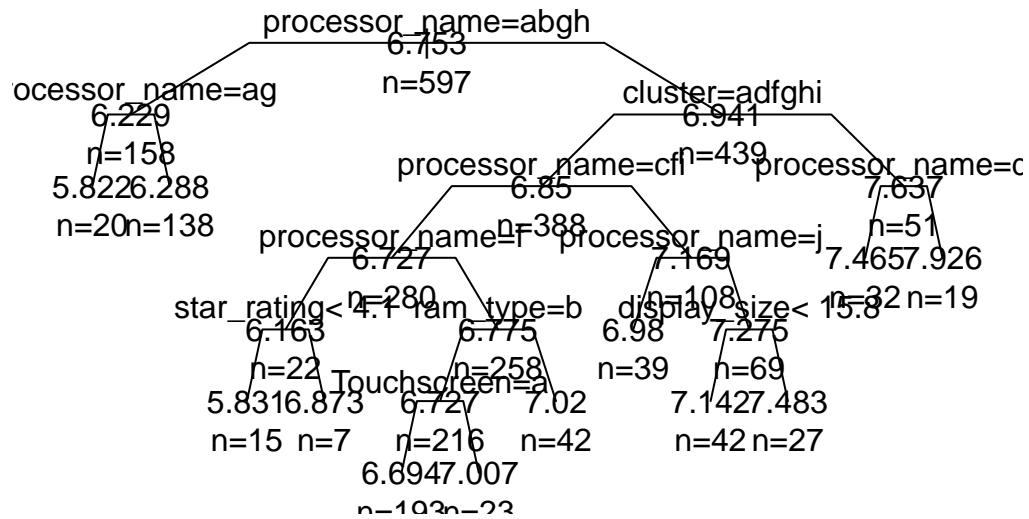
Decision Tree

We wanted further visualization of the laptop price variable, hence we made a regression tree to visualize this.

```
tree.out.1 <- rpart(log_latest_price ~ ., data = pred.laptop.train,
                   parms = list(split="information"),
                   control = rpart.control(minsplit=20))
```

```
#Create a plot of the classification tree.
#Code to plot the tree.
plot(tree.out.1, uniform=TRUE, branch=0.6, margin=0.05)
text(tree.out.1, all=TRUE, use.n=TRUE)
title("Laptop Price Regression Tree")
```

Laptop Price Regression Tree



Looking at the results we were able to see that we got an R-squared value of around .68 for this model

Random Forest Model

We then used the random forest model to see if this would be a better model, and in fact was.

```
laptop.train.rf <- randomForest(log_latest_price ~ .,
                                data = pred.laptop.train,
                                importance=TRUE)


print(laptop.train.rf$importance)
```

##	%IncMSE	IncNodePurity
## brand	0.0059684506	3.5434391
## processor_brand	0.0153051988	2.9744342
## processor_name	0.0923899754	28.9346092
## ram_type	0.0087423660	4.3039938
## os	0.0216309357	5.7705393
## os_bit	0.0006633697	0.8664901
## graphic_card_gb	0.0204087236	7.5760371
## weight	0.0042694026	1.4592317
## display_size	0.0275332441	11.0879834
## warranty	0.0036213301	1.9348522
## Touchscreen	0.0095114446	3.3301613
## msoffice	0.0028293931	1.1686366


```
## star_rating      0.0149630137      6.4254561
## ratings          0.0206146725      6.7990646
## reviews         0.0157708219      5.9301584
## display_size_given 0.0064951843      1.7197173
## ram_gb_cat       0.0225982440     12.6229860
## ssd_cat          0.0214271808     10.3615372
## hdd_cat          0.0017975182      0.7782949
## cluster          0.0872089169     37.2697937
```

```
varImpPlot(laptop.train.rf)
```

laptop.train.rf



```
optimum <- which.max(laptop.train.rf$importance[, "%IncMSE"])
opt.var <- laptop.train.rf$importance[optimum, 0, drop=FALSE]

print("The most predictive variable with regard to price is:")
```

```
## [1] "The most predictive variable with regard to price is:"
```

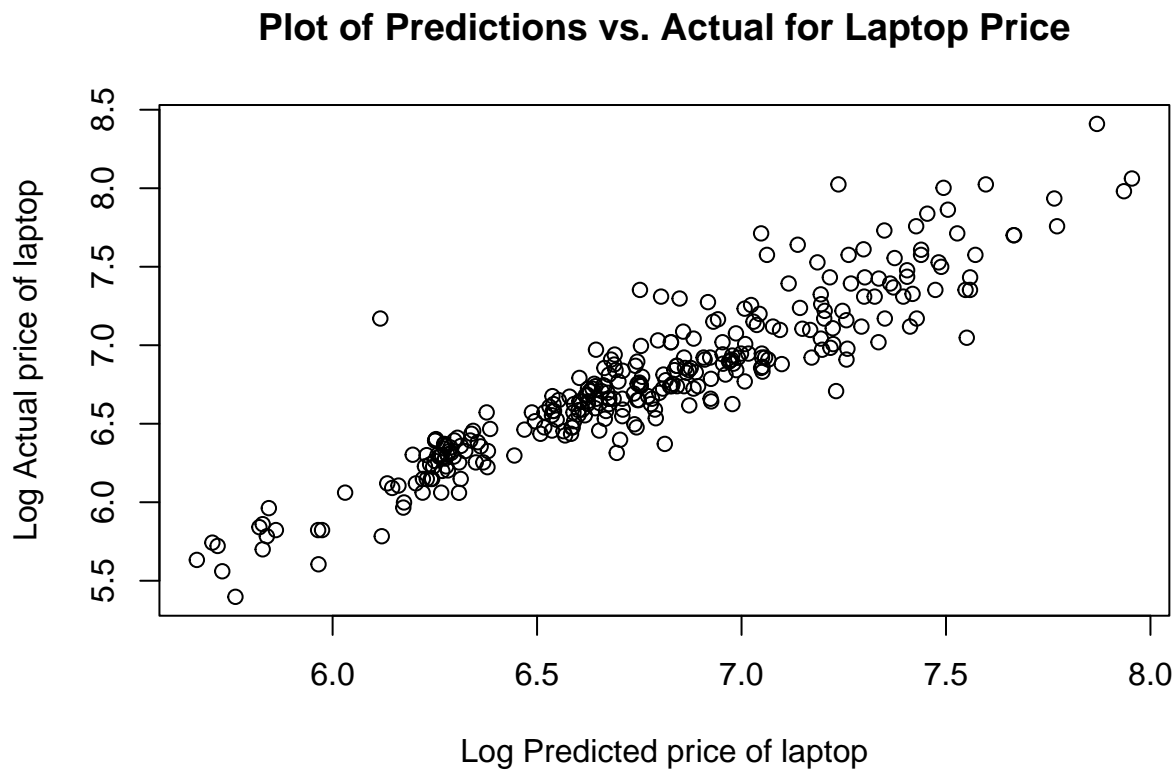
```
print(opt.var)
```

```
##
## processor_name
```

```
val.preds.rf <- predict(laptop.train.rf, # The forest
                        newdata = pred.laptop.validation, # The values of x to do prediction at
                        type = c("response")
                        )

# Code to plot the predictions against the actual values

plot(val.preds.rf, pred.laptop.validation$log_latest_price,
     main = "Plot of Predictions vs. Actual for Laptop Price",
     xlab = "Log Predicted price of laptop",
     ylab = "Log Actual price of laptop")
```



Looking at the summary of the model we were able to get an R-squared value of .78

Best Fitting Model

Through the models created it was clear that the Random Forest provided us with the best analysis of the laptop price data. This was seen as it had the highest R-squared value of around .78. It was seen that processor name variables followed by cluster and ssd_cat were the most predictive towards laptop price.

Conclusion

Conclusion

This data set was interesting to visualize and to implement a clustering methodology seemed to be more practice than using cross folding. Seeing our best fitting model was the random forest model, was great, however, was still not as accurate we would have liked it to be. Being able to model this data set and see some correlation of variables to predict laptop price could be used for consumers to have a second thought of which laptop brand they should consider purchasing. In the future if we wanted to expand on the data set, we could instead try using cross folding to model the data, as well as using an XGboost model to look further at the prediction of laptop price.